


GT: ALS
F: rust
G: ALS




GT: ALS
F: rust
G: ALS



Low-Light Qualitative Results

Examples from the low-light corrupted Beans test split. We show cases where the fixed recipe fails and the guided recipe remains correct.

GT: ALS
F: healthy
G: ALS



GT: rust
F: ALS
G: rust



Real-data benchmark

```
train = 1034 images  
val = 133 images  
test = 128 images  
train device = mps  
eval = clean + low-light test
```

Beans Robustness

